

UIMA Components

Input

- OpenTrivaQA parser
 - Parses files in OpenTriviaQA format
 - Each question is a document
 - For each question the following annotations are added
 - Question → the question text
 - Answer → the correct answer
 - Answer candidates → all four answer candidates; each candidate has its own annotation

Kommentiert [TK1]: This may be dropped because I don't see a use case in our system for the answer candidates...

Question Processing

- Keyword extraction
 - The question should be analyzed by other components (e.g. POS tagger) before
 - Extract all key words from question
 - Beyond the baseline: the component may add an importance score to each keyword

Tag Cloud Enrichment

- Category and hypernym detection
 - This is just a single component which is internally multi-threaded
 - For each resource (e.g. Wikipedia, WordNet) a separate thread is started
 - Each resource creates a list of categories/hypernyms
 - After the research for each resource is finished, the list is set as the corresponding feature in the answer annotation (thereby building the "large tag cloud")
 - **Parameters**
 - Annotation type: over which annotation to iterate; this will be used to run the analysis engine on the correct answer first and later on the answer candidates

Kommentiert [TK2]: As of my research, it seems to be non-trivial to parallelize UIMA analysis engines

Kommentiert [TK3]: That's why there are the two distinct types "correct answer" and "candidate answer"

Candidate Extraction

- Category ranking
 - Iterates over all categories and hypernyms (possibly in parallel) of the correct answer
 - Uses a measurement to derive a score for each category which reflect its relevance
 - For instance, page rank could be used for Wikipedia categories
 - The measurement may change to allow different system configurations
 - The scores are assigned to the categories
 - **Parameters**

- Relevance measure
- Candidate extraction
 - Extract categories/hypernyms from the correct answer with the highest relevance score
 - For each of the extracted categories, retrieve candidates
 - Create an “answer candidate” annotation for each retrieved candidate
 - **Parameters**
 - How many categories/hypernyms to use
- Synonym resolution
 - Iterate over all “answer candidate” annotations
 - Perform synonym resolution with the correct answer
 - E.g. by using WordNet
 - Remove all synonymous candidate answers by removing the corresponding “answer candidate” from the index

Kommentiert [TK4]: This is probably not easy

Kommentiert [TK5]: Maybe add the correct answer's keywords to it?

Similarity Detection

- Similarity detection
 - Iterate over all “answer candidate” annotations
 - Perform the similarity measure between the candidate's tag cloud and the correct answer's tag cloud
 - Assign the resulting scores to the answer candidates
 - **Parameters**
 - The similarity measure

Kommentiert [TK6]: Maybe perform the similarity calculation in a separate thread for each answer candidate

CAS Consumer

- Candidate selection
 - Iterate over all “answer candidate” annotations
 - Select the candidates with the highest scores and output them
 - **Parameters**
 - Number of candidates to select

Kommentiert [TK7]: This may be used as an input for another system

UIMA Types

Question

- Standard UIMA annotation
- No special features

Answer

- Annotation for an answer (correct or candidate)
- Features
 - Keywords
 - Categories
 - Hypernyms
 - Answer type

Kommentiert [TK8]: All of them implement a common interface which allows unified access to them

Kommentiert [TK9]: This, however, would require an additional analysis engine

Correct Answer

- Inherits everything from the answer type
- Is used to distinguish the correct answer annotation from the candidate annotation

Answer Candidate

- Inherits everything from the answer type
- Features
 - Similarity score
- Is used for iterating over answer candidates

Final Pipeline

